

Store trip analysis

by
Rui Furtado

Data analysis

Libraries

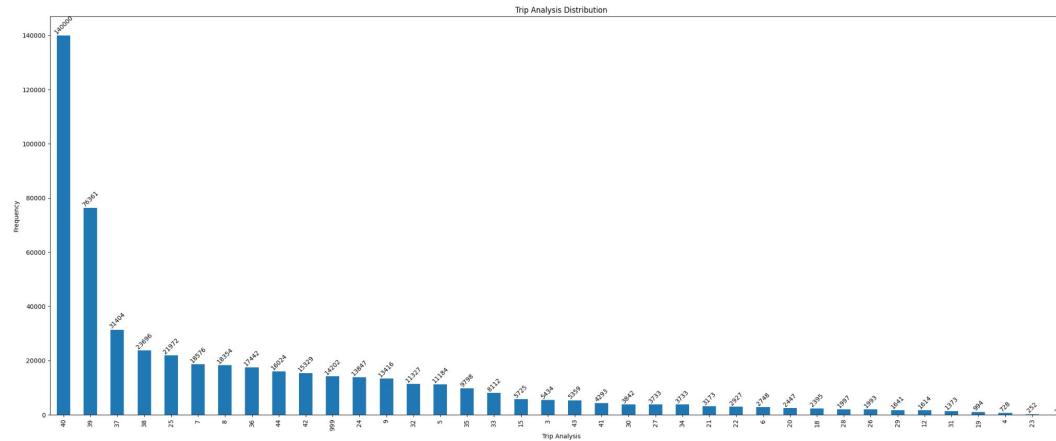
For developing this project the following libs were installed:

- [fastparquet](#): parquet engine
- [fastpivot](#): “shady lib” that creates a fast implementation of the pandas function pivot_table using scipy sparse matrices. It was great!
- [matplotlib](#): for simple plotting
- [pandas](#): data preprocessing
- [pyarrow](#): apache arrow engine for saving parquet train data
- [scikit-learn](#): ML algorithms
- [uv](#): package manager

To activate the env use uv sync followed by source .venv/bin/activate or just pip install -r requirements if you don't want to use a dedicated environment. The project uses python 3.12

TripType

- Target column. looks pretty unbalanced. We will need to tackle this issue when creating our prediction algorithm either with:
 - using algorithms that deal with this issue
 - over/under sample in order to have more/less examples
 - 38 classes represented



VisitNumber

- Will be used as primary key
- 76605 elements
- No nulls
- Will be used as index for our data
- Multiple records by visit
- model will predict by store trip so records need to be aggregated by VisitNumber

Weekday

- String with weekday of record
- No nulls
- Values need to be correctly standardized - lower, trimmed
- Cyclical encoding was chosen
- Features can be derived from it - `is_weekend`

```
Weekday
Sunday      104555
Saturday    89807
Friday       72054
Monday      62828
Wednesday   57141
Tuesday     55902
Thursday    48186
SATURDAY    7000
THURSDAY    6000
FrIdAy      5000
MonDAY      4000
sunday      3000
tuesday     2000
Name: count, dtype: int64
```

Upc

- Product code
- Number of nans - 3331 - will be dropped
- High cardinality - 85895 elements
- No nulls
- Will be one-hot encoded to generate a column by product (aggregated by visit)

ScanCount

- Product code
- High cardinality - 85895 elements
- No nulls
- Will be one-hot encoded to generate a column by product (aggregated by visit)
- Presents negative numbers (from -1 to -10). These were kept since they were returns
- Used to create total count of products (after agg)

DepartmentDescription

- Product categorization describe department
- 99 elements
- Should be cleaned - lowercase, trim, space removal
- After cleaning we got 67 elements
- Will be also one-hot encoded to generate 1 column by department (agg by VisitNumber)

FineLineNumber

- Another categorization of the product
- Presented 3331 nulls
- Was dropped since we wanted to use UPC for column dummies
- Should have used this one since cardinality was lower (5126 elements vs 85895 elements)

Preprocessing

Encoding of text categories

- Used regex for removing non-chars, trimmed and lowered
- Applied to weekday and department description

```
# Function for standardization
def standardize_text_categories(cat:str) -> str:
    """Standardizes given text

    Args:
        cat (str): category to be standardize

    Returns:
        standardize text
    """
    return re.sub(r'[^a-zA-Z]+', '', cat.strip().lower()) # remove non-letters
```

Aggregation by visit and total

- Aggregation by visit number and cols that identify a visit - groupby VisitNumber, TripType, Weekday
- Sum of ScanCount to create Total - total number of products acquired per visit
- Records reduced from 517475 to 75456

```
# Remove rows with no Upc
data_cleaned = data_cleaned.dropna(subset=['Upc'])

visit_agg = data_cleaned[['VisitNumber', 'TripType', 'Weekday', 'ScanCount']].groupby(['VisitNumber', 'TripType', 'Weekday'], as_index=False).sum()
visit_agg=visit_agg.set_index('VisitNumber').rename(columns={"ScanCount":"Total"}) # set index to VisitNumber to be able to join and rename ScanCount
visit_agg
```

Pivot from UPC, visitNumber and ScanCount

- Creation of Pivot table for getting a dataframe with visit numbers by number of products acquired by product type
- Usage of fastpivot lib to address memory issues caused by high cardinality - sparse matrix creation
- Inner join with aggregated visit table by index (VisitNumber) - 74566 x 85894 matrix

The screenshot shows a Jupyter Notebook cell with the following content:

```
from fastpivot import pivot_sparse
pivot_upc=pivot_sparse(data_cleaned, index='VisitNumber', columns='Upc', values='ScanCount', fill_value=0).fillna(0)
pivot_upc
```

The output of the code is a sparse matrix visualization. The columns are labeled with UPC codes (e.g., 8.340000e+02, 3.032000e+03, etc.) and the rows are labeled with VisitNumbers (e.g., 8, 11, 12, 15, 19, ..., 191347). The matrix is mostly filled with zeros, indicating that most visit numbers do not correspond to specific UPC codes. The matrix has 75456 rows and 85894 columns.

75456 rows x 85894 columns

Departement description dummies

- One-hot encoding of department description and aggregation by visit number
- Group by VisiNumber to get frequency of departments by visit
- Usage of sparse=True in pd.get_dummies to use sparse matrix
- Joined with previous aggregated matrix by VisitNumber

```
department_dummies = pd.get_dummies(
    data_cleaned.set_index('VisitNumber').DepartmentDescription, dtype=int, sparse=True
).groupby('VisitNumber').sum()
department_dummies
✓ 1m 27.0s
```

VisitNumber	accessories	automotive	bakery	bathandshower	beauty	bedding	booksandmagazines	boyswear	brasshapewear	camerasandsupplies	...	seafood	seasonal	servicedeli
8	0	0	0	0	0	0	0	0	0	0	...	0	0	0
11	0	0	0	0	0	0	0	0	0	0	...	0	0	0
12	0	0	0	0	0	0	0	2	0	0	...	0	0	0
15	0	0	0	0	0	0	0	0	0	0	...	0	0	0
19	1	0	0	0	0	0	0	0	0	0	...	0	0	0
...
191335	0	0	0	0	0	0	0	0	0	0	...	0	0	0
191343	0	0	0	0	0	0	0	0	0	0	...	0	0	0
191344	0	0	0	0	4	0	0	0	0	0	...	0	0	0
191346	0	0	0	0	0	0	0	0	0	0	...	0	0	0
191347	0	0	0	0	0	0	0	0	0	0	...	0	0	0

6 rows × 67 columns

Weekday encoding

- Cyclical encoding with cosine and sine to preserve proximity of week final with week beginning
- Creation of is_weekend
- Creation of final dataframe - 75456 x 85966

```
DAY_MAPPING= {
    'monday': 1, 'tuesday': 2, 'wednesday': 3, 'thursday': 4,
    'friday': 5, 'saturday': 6, 'sunday': 7
}

day_num = agg['Weekday'].map(DAY_MAPPING)

# Encode weekday into cyclical features
agg['weekday_sin'] = np.sin((2 * np.pi / 7)* day_num)
agg['weekday_cos'] = np.cos((2 * np.pi / 7)* day_num)

# check when we have a weekend
agg['is_weekend'] = agg['Weekday'].isin(['saturday', 'sunday']).astype(int)
agg = agg.drop('Weekday', axis=1)
agg

✓ 0.4s
```

Data storage

- Data stored in parquet - to_csv too slow for high volume of data
- Needed to convert pandas types from sparse to int and float
- High memory usage! Should have used FineLineNumber instead of Upc for one-hot encoding

```
type_dict = {**dict.fromkeys(agg.columns, int), "weekday_sin": float, "weekday_cos": float}
agg.info(memory_usage='deep')
agg = agg.astype(type_dict)
agg.info(memory_usage='deep')

✓ 8.2s

<class 'pandas.DataFrame'>
Index: 75456 entries, 8 to 191347
Columns: 85966 entries, TripType to is_weekend
dtypes: Sparse[int64, 0](85961), float64(2), int64(3)
memory usage: 12.1 MB
<class 'pandas.DataFrame'>
Index: 75456 entries, 8 to 191347
Columns: 85966 entries, TripType to is_weekend
dtypes: float64(2), int64(85964)
memory usage: 48.3 GB

agg.to_parquet('processed_data.parquet')
```

Modeling

Data splitting

- Data was split into train and test - 80% train and 20% for test
- Stratified to have a test set that resembles class distribution of train - since we have imbalance of classes this is needed in order to have all classes present

```
from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=4,
    stratify=y,
    shuffle=True
)
✓ 1m 22.1s
```

Baseline model

- Random forest model was chosen for baseline prediction:
 - No need for feature scaling
 - Resistant to overfitting
 - Only looks at a subset of features at a time
 - Handles non-linear relationships
 - Provides feature importance
 - Provides balancing option off the shelf
- Model took around 40 min to train due to high cardinality ! (Mistake - trained with 500 trees and full data. train and test results were too similar. After correction 9min (200 trees and corrected data))
- Model was immediately saved to pkl file after training to avoid re-fitting

Baseline model

- Parameters chosen
 - 200 trees
 - Maximum depth of each tree 25 branches to try to prevent overfitting
 - class_weight=balanced to penalize rare class misclassification

```
from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(
    n_estimators=200,
    max_depth=20,
    random_state=43,
    class_weight='balanced',
    n_jobs=-1,
).fit(X_train.values,y_train.values)

✓ 9m 6.8s
```

Results train prediction

- Results were not great even for prediction with X_train
 - Model is underfitting
 - Some classes still present good results
 - F1-score of 51% - same as tossing a coin
 - More predictive power maybe would be helpful

	precision	recall	f1-score	support					
3.0	0.94	0.93	0.94	2327					
4.0	0.10	0.96	0.19	226					
5.0	0.88	0.14	0.24	2181					
6.0	0.40	0.93	0.56	808					
7.0	0.63	0.57	0.60	3694					
8.0	0.82	0.36	0.50	7826					
9.0	0.97	0.01	0.02	6028					
12.0	0.67	0.91	0.77	163					
14.0	0.05	1.00	0.10	3					
15.0	0.39	0.58	0.47	622					
18.0	0.32	0.85	0.47	354					
19.0	0.24	0.97	0.39	245					
20.0	0.29	0.99	0.45	407					
21.0	0.30	0.90	0.45	418					
22.0	0.61	0.35	0.45	600					
23.0	0.14	1.00	0.25	84					
24.0	0.49	0.74	0.59	1662					
25.0	0.49	0.71	0.58	2374					
26.0	0.25	0.94	0.39	317					
27.0	0.36	0.93	0.52	508					
28.0	0.27	0.99	0.42	295					
29.0	0.97	0.41	0.58	277					
30.0	0.32	0.84	0.46	682					
31.0	0.56	1.00	0.72	369					
32.0	0.51	0.93	0.66	1280					
33.0	0.38	0.74	0.50	858	accuracy				
34.0	0.33	0.96	0.49	454	macro avg	0.54	0.69	0.48	60364
35.0	0.42	0.47	0.44	1274	weighted avg	0.72	0.51	0.47	60364

Results test prediction

- Results were also not good for testset prediction
- F1-score of 46%

	precision	recall	f1-score	support					
3.0	0.95	0.92	0.93	582					
4.0	0.09	0.80	0.16	56					
5.0	0.65	0.20	0.30	545					
6.0	0.33	0.90	0.49	202					
7.0	0.60	0.56	0.58	924					
8.0	0.78	0.35	0.49	1957					
9.0	0.94	0.09	0.16	1507					
12.0	0.20	0.38	0.26	40					
14.0	0.00	0.00	0.00	0					
15.0	0.26	0.54	0.35	156					
18.0	0.23	0.65	0.34	88					
19.0	0.25	0.90	0.39	61					
20.0	0.29	0.85	0.43	102					
21.0	0.26	0.78	0.39	105					
22.0	0.41	0.29	0.34	150	35.0	0.34	0.35	0.35	319
23.0	0.17	0.95	0.29	21	36.0	0.44	0.40	0.42	479
24.0	0.47	0.64	0.55	415	37.0	0.30	0.77	0.43	449
25.0	0.52	0.67	0.59	594	38.0	0.23	0.39	0.29	463
26.0	0.20	0.95	0.32	79	39.0	0.53	0.01	0.01	1586
27.0	0.33	0.80	0.47	127	40.0	0.51	0.77	0.61	993
28.0	0.21	0.84	0.33	74	41.0	0.00	0.00	0.00	92
29.0	0.32	0.10	0.15	69	42.0	0.46	0.02	0.04	292
30.0	0.28	0.55	0.37	170	43.0	0.00	0.00	0.00	145
31.0	0.53	0.98	0.68	92	44.0	0.37	0.07	0.12	187
32.0	0.50	0.87	0.63	320	999.0	1.00	0.79	0.88	1322
33.0	0.26	0.74	0.38	215	accuracy			0.46	15092
34.0	0.30	0.89	0.45	114	macro avg	0.38	0.55	0.37	15092
					weighted avg	0.60	0.46	0.43	15092

Model #2 No Upc

- Since training took some time, I decided based on RF feature importance, to remove all the features related with Upc

	Feature	Importance
0	Total	0.028035
85945	petsandsupplies	0.022452
85948	playersandelectronics	0.017128
85961	wireless	0.016792
85896	automotive	0.016388
85912	electronics	0.015778
85951	produce	0.015512
85958	sportinggoods	0.015224
85914	financialservices	0.014589
85929	infantconsumablehardlines	0.014288
85773	692303000000.0	0.013809
85934	lawnandgarden	0.013122
60699	68113102889.0	0.012862
85927	impulsemerchandise	0.012508
85960	toys	0.012196
85913	fabricsandcrafts	0.012036
85939	officesupplies	0.011325
57899	64541689243.0	0.011053
85906	celebration	0.010997
85956	shoes	0.010514

Model #2 No Upc - Results train prediction

- Results improved 91% for train
- Model had the same

params but max_depth was set
to None

	precision	recall	f1-score	support					
3.0	0.81	0.98	0.89	2158					
4.0	0.25	0.94	0.40	203					
5.0	0.91	0.74	0.81	2067					
6.0	0.84	0.97	0.90	750					
7.0	0.91	0.95	0.93	3449					
8.0	0.91	0.86	0.89	7364					
9.0	0.95	0.66	0.78	5684					
12.0	0.82	0.99	0.90	164					
14.0	1.00	1.00	1.00	2					
15.0	0.90	0.95	0.92	594					
18.0	0.77	0.94	0.84	332					
19.0	0.41	0.99	0.58	233					
20.0	0.72	0.98	0.83	386					
21.0	0.88	0.98	0.93	373					
22.0	0.84	0.78	0.81	539					
23.0	0.42	1.00	0.59	75					
24.0	0.90	0.98	0.94	1537					
25.0	0.96	0.98	0.97	2246					
26.0	0.73	0.99	0.84	289					
27.0	0.87	0.98	0.92	481					
28.0	0.64	0.98	0.78	278					
29.0	0.60	0.94	0.74	256					
30.0	0.68	0.99	0.81	650					
31.0	0.68	0.99	0.81	341					
32.0	0.89	1.00	0.94	1185					
33.0	0.98	0.99	0.99	820					
34.0	0.88	0.99	0.93	408					
			accuracy			0.91	56592		
			macro avg			0.84	0.95	0.88	56592
			weighted avg			0.93	0.91	0.91	56592

Model #2 No Upc - Results test prediction

- Test set results improved to 64% F1-score
- Still with problems in

minority classes

	precision	recall	f1-score	support	
3.0	0.78	0.98	0.87	751	
4.0	0.17	0.49	0.25	79	
5.0	0.58	0.53	0.55	659	
6.0	0.66	0.74	0.70	260	
7.0	0.67	0.65	0.66	1169	
8.0	0.78	0.77	0.77	2419	
9.0	0.72	0.52	0.60	1851	
12.0	0.07	0.05	0.06	39	
14.0	0.00	0.00	0.00	1	
15.0	0.46	0.48	0.47	184	
18.0	0.29	0.47	0.36	110	
19.0	0.20	0.55	0.30	73	
20.0	0.46	0.77	0.58	123	
21.0	0.56	0.68	0.62	150	
22.0	0.42	0.35	0.38	211	
23.0	0.32	0.50	0.39	30	
24.0	0.57	0.49	0.53	540	
25.0	0.64	0.69	0.66	722	
26.0	0.38	0.47	0.42	107	
27.0	0.46	0.65	0.54	154	
28.0	0.33	0.62	0.43	91	
29.0	0.14	0.17	0.15	90	
30.0	0.39	0.56	0.46	202	
31.0	0.52	0.90	0.66	120	
32.0	0.65	0.85	0.73	415	
33.0	0.62	0.58	0.60	253	
34.0	0.57	0.61	0.59	160	
					accuracy
					macro avg
					weighted avg
					0.64
					18864
					0.52
					0.48
					18864
					0.64
					18864
					0.63
					18864

Model #3 SMOTE - Results train prediction

- For model 3 we decided to use SMOTE with no Upc to generate more samples and try to overcome issues with minority classes
- Max_depth was set to

30 this time

	precision	recall	f1-score	support					
3.0	0.91	0.96	0.93	7364					
4.0	0.76	1.00	0.86	7364					
5.0	0.91	0.65	0.76	7364					
6.0	0.95	0.96	0.95	7364					
7.0	0.90	0.76	0.83	7364					
8.0	0.56	0.83	0.67	7364					
9.0	0.88	0.41	0.56	7364					
12.0	1.00	0.99	0.99	7364					
14.0	1.00	1.00	1.00	7364					
15.0	0.95	0.82	0.88	7364					
18.0	0.93	0.90	0.92	7364					
19.0	0.80	1.00	0.89	7364					
20.0	0.95	1.00	0.97	7364					
21.0	0.96	0.97	0.97	7364					
22.0	0.97	0.78	0.87	7364					
23.0	0.98	1.00	0.99	7364					
24.0	0.87	0.94	0.90	7364					
25.0	0.89	0.88	0.89	7364					
26.0	0.93	1.00	0.96	7364					
27.0	0.97	1.00	0.98	7364					
28.0	0.91	1.00	0.95	7364					
29.0	0.95	0.92	0.93	7364					
30.0	0.89	0.96	0.92	7364					
31.0	0.96	0.96	0.96	7364					
32.0	0.95	0.98	0.96	7364					
33.0	0.94	0.97	0.95	7364					
34.0	0.96	1.00	0.98	7364					
35.0	0.74	0.90	0.81	7364					
					accuracy				0.91
					macro avg	0.92	0.91	0.91	279832
					weighted avg	0.92	0.91	0.91	279832

Model #3 SMOTE - Results test prediction

- Despite balancing classes result did not improve
- F1-score 61%

	precision	recall	f1-score	support					
3.0	0.78	0.96	0.86	751					
4.0	0.16	0.59	0.25	79					
5.0	0.54	0.50	0.52	659					
6.0	0.56	0.82	0.67	260					
7.0	0.74	0.57	0.65	1169					
8.0	0.74	0.80	0.77	2419					
9.0	0.82	0.33	0.47	1851					
12.0	0.10	0.05	0.07	39					
14.0	0.00	0.00	0.00	1					
15.0	0.48	0.52	0.50	184					
18.0	0.29	0.64	0.40	110	35.0	0.47	0.62	0.53	388
19.0	0.22	0.70	0.33	73	36.0	0.57	0.65	0.61	576
20.0	0.35	0.91	0.51	123	37.0	0.53	0.57	0.55	538
21.0	0.50	0.74	0.60	150	38.0	0.40	0.45	0.42	612
22.0	0.41	0.37	0.39	211	39.0	0.60	0.49	0.54	1965
23.0	0.33	0.60	0.42	30	40.0	0.78	0.91	0.84	1290
24.0	0.57	0.53	0.55	540	41.0	0.33	0.05	0.09	121
25.0	0.63	0.62	0.62	722	42.0	0.42	0.17	0.24	371
26.0	0.31	0.67	0.43	107	43.0	0.25	0.03	0.06	154
27.0	0.44	0.75	0.55	154	44.0	0.33	0.07	0.11	230
28.0	0.27	0.82	0.41	91	999.0	1.00	0.71	0.83	1656
29.0	0.15	0.16	0.15	90	accuracy				0.61 18864
30.0	0.33	0.66	0.44	202	macro avg				0.47 18864
31.0	0.53	0.92	0.67	120	weighted avg				0.61 18864
32.0	0.59	0.87	0.71	415					
33.0	0.49	0.68	0.57	253					
34.0	0.44	0.78	0.56	160					

Model #4 K best

- Based on the previous features importance lets explore using only a predefined number of features k=1000. This will include Upc and the rest

	Feature	Importance
0	Total	0.028035
85945	petsandsupplies	0.022452
85948	playersandelectronics	0.017128
85961	wireless	0.016792
85896	automotive	0.016388
85912	electronics	0.015778
85951	produce	0.015512
85958	sportinggoods	0.015224
85914	financialservices	0.014589
85929	infantconsumablehardlines	0.014288
85773	692303000000.0	0.013809
85934	lawnandgarden	0.013122
60699	68113102889.0	0.012862
85927	impulsemerchandise	0.012508
85960	toys	0.012196
85913	fabricsandcrafts	0.012036
85939	officesupplies	0.011325
57899	64541689243.0	0.011053
85906	celebration	0.010997
85956	shoes	0.010514

Model #4 K best - Results train prediction

- Training improved a bit 94%. This might represent overfitting

	precision	recall	f1-score	support					
3.0	0.97	1.00	0.98	2180	35.0	0.98	0.99	0.98	1202
4.0	0.41	0.87	0.56	210	36.0	0.97	0.97	0.97	1781
5.0	0.90	0.87	0.88	2032	37.0	1.00	1.00	1.00	1702
6.0	0.88	0.97	0.92	733	38.0	0.99	0.99	0.99	1744
7.0	0.94	0.97	0.96	3489	39.0	1.00	1.00	1.00	5943
8.0	0.93	0.91	0.92	7414	40.0	1.00	1.00	1.00	3685
9.0	0.96	0.71	0.82	5674	41.0	0.99	1.00	1.00	351
12.0	0.99	1.00	1.00	152	42.0	1.00	0.99	1.00	1081
14.0	1.00	1.00	1.00	2	43.0	1.00	1.00	1.00	530
15.0	0.94	0.96	0.95	572	44.0	1.00	1.00	1.00	720
18.0	0.82	0.96	0.88	341	999.0	0.99	0.92	0.95	4911
19.0	0.62	0.93	0.74	222	accuracy			0.94	56592
20.0	0.83	0.97	0.89	385	macro avg	0.88	0.97	0.92	56592
21.0	0.90	0.99	0.95	400	weighted avg	0.95	0.94	0.94	56592
22.0	0.74	0.91	0.82	572					
23.0	0.65	1.00	0.79	77					
24.0	0.89	0.99	0.94	1581					
25.0	0.96	0.99	0.97	2204					
26.0	0.76	0.99	0.86	292					
27.0	0.90	0.99	0.94	466					
28.0	0.66	1.00	0.80	275					
29.0	0.67	0.97	0.79	262					
30.0	0.68	0.99	0.81	618					
31.0	0.77	1.00	0.87	344					
32.0	0.90	1.00	0.95	1216					
33.0	0.99	0.99	0.99	813					
34.0	0.90	0.99	0.94	416					

Model #4 K best - Results test prediction

- Results slightly improve but we still have the same issues as before
- F1-score of 65%, maybe due to better score on test?
- Would need to cross-validate to better understand

	precision	recall	f1-score	support					
3.0	0.94	0.98	0.96	729					
4.0	0.17	0.26	0.21	72					
5.0	0.63	0.64	0.63	694					
6.0	0.72	0.70	0.71	277					
7.0	0.66	0.69	0.68	1129					
8.0	0.79	0.79	0.79	2369					
9.0	0.69	0.56	0.62	1861					
12.0	0.22	0.04	0.07	51					
14.0	0.00	0.00	0.00	1					
15.0	0.51	0.37	0.43	206					
18.0	0.32	0.50	0.39	101					
19.0	0.29	0.33	0.31	84					
20.0	0.49	0.66	0.56	124					
21.0	0.55	0.63	0.58	123					
22.0	0.39	0.53	0.45	178					
23.0	0.35	0.32	0.33	28					
24.0	0.56	0.60	0.58	496					
25.0	0.66	0.68	0.67	764					
26.0	0.36	0.41	0.39	104					
27.0	0.49	0.73	0.59	169					
28.0	0.30	0.44	0.36	94					
29.0	0.23	0.23	0.23	84					
30.0	0.38	0.44	0.41	234					
31.0	0.60	0.89	0.72	117					
32.0	0.60	0.79	0.68	384					
33.0	0.60	0.52	0.56	260					
34.0	0.58	0.62	0.60	152					
					accuracy				0.65
					macro avg	0.51	0.50	0.48	18864
					weighted avg	0.66	0.65	0.64	18864

TODOS

This challenge is not completed and due to time constraints I was not able to fully implement all the strategies that I had in my mind for reaching a better outcome. Some steps I would take to improve the presented results:

- Explore the usage of FineLineNumber instead of Upc. Upc columns could be discarded since they were barely used by the final model
- Employ a cross validation strategy to access results- Current results are not 100% trustworthy since this was not applied. The stratified version should be used since this is a multiclass problem.
- Use a gridsearch to explore better hyperparameters for the random forest.
- Explore other models such as GradientBoostingTress or XGboost
- Explore other strategies of feature selection for example having k-best in GridSearch
- Explore using bins for some of the features - for example aggregating similar departments or even removing or join of TripTypes